

## hw5

October 2, 2025

**Problem 1 :** Using the images (aerobic-[001-022].bmp) provided on the class materials site, experiment with simple “motion detection” between consecutive frames using (abs) image differencing. Clean-up and remove any tiny regions (e.g., use techniques such as bwareaopen, median filtering, etc.). Experiment with different thresholds. [2 pts]

```
[ ]: # Q2: MEI & MHI
import os, glob
import numpy as np
import matplotlib.pyplot as plt
from PIL import Image

DATA_DIR = "data/"
PATTERN = "aerobic-*.bmp"

# Motion mask params (same as Q1)
T = 0.05 # fixed threshold in [0,1]
MED_K = 3 # median filter window (odd)
MIN_SIZE = 20 # min component size (px)

# MHI params for this sequence (i goes 2..22 -> 21 diffs total)
TAU = 21 # decay horizon in steps
DELTA = 1.0 / TAU # per-step decay in normalized [0,1]
```

```
[22]: def load_gray01(path):
        return np.asarray(Image.open(path).convert("L"), np.float32) / 255.0

def load_sequence(data_dir=DATA_DIR, pattern=PATTERN):
    paths = sorted(glob.glob(os.path.join(data_dir, pattern)))
    if not paths: raise FileNotFoundError("No BMPs found.")
    frames = [load_gray01(p) for p in paths]
    return frames, paths

def pad_reflect(a, r):
    return np.pad(a, r, mode="reflect")

def median_filter2d(img, k=3):
    # simple, pure-NumPy median filter (k odd)
    assert k % 2 == 1
```

```

r = k // 2
P = pad_reflect(img, r)
H, W = img.shape
out = np.empty_like(img)
for y in range(H):
    rows = P[y:y+k]
    win = np.stack([rows[:, x:x+W] for x in range(k)], axis=0) # k x k x W
    out[y] = np.median(win.reshape(k*k, W), axis=0)
return out

def remove_small_regions(bw, min_size=MIN_SIZE):
    # 8-connected flood fill; bw is {0,1} or bool
    bw = (bw > 0).astype(np.uint8)
    H, W = bw.shape
    lab = np.zeros_like(bw, np.int32)
    cur = 0
    nbrs = [(-1,-1), (-1,0), (-1,1), (0,-1), (0,1), (1,-1), (1,0), (1,1)]
    from collections import deque
    for y in range(H):
        for x in range(W):
            if bw[y,x] and lab[y,x]==0:
                cur += 1; q = deque([(y,x)]); lab[y,x]=cur; size=0
                while q:
                    cy,cx=q.popleft(); size+=1
                    for dy,dx in nbrs:
                        ny,nx=cy+dy,cx+dx
                        if 0<=ny<H and 0<=nx<W and bw[ny,nx] and lab[ny,nx]==0:
                            lab[ny,nx]=cur; q.append((ny,nx))
                if size < min_size: lab[lab==cur]=0
    return (lab>0).astype(np.uint8)

# ---- shape moments (raw/central/normalized + Hu) ----
def spatial_moments(img):
    """
    img: 2D array in [0,1] (float) or {0,1} (binary). Treated as intensity (MEI/
    ↪ MHI).
    Returns dict with m_ij, mu_ij, eta_ij and 7 Hu invariants.
    """
    I = img.astype(np.float64)
    H, W = I.shape
    yy, xx = np.mgrid[0:H, 0:W]

    m00 = I.sum()
    # guard against empty images
    if m00 <= 1e-12:
        return dict(m00=0, cx=np.nan, cy=np.nan, hu=[np.nan]*7)

```

```

m10 = (xx*I).sum(); m01 = (yy*I).sum()
cx, cy = m10/m00, m01/m00

x = xx - cx; y = yy - cy
mu20 = (x**2 * I).sum(); mu02 = (y**2 * I).sum(); mu11 = (x*y * I).sum()
mu30 = (x**3 * I).sum(); mu03 = (y**3 * I).sum()
mu21 = (x**2 * y * I).sum(); mu12 = (x * y**2 * I).sum()

# normalized central moments (scale invariance)
eta20 = mu20 / (m00**2)
eta02 = mu02 / (m00**2)
eta11 = mu11 / (m00**2)
eta30 = mu30 / (m00**2.5)
eta03 = mu03 / (m00**2.5)
eta21 = mu21 / (m00**2.5)
eta12 = mu12 / (m00**2.5)

# Hu invariants (1..7)
phi1 = eta20 + eta02
phi2 = (eta20 - eta02)**2 + 4*eta11**2
phi3 = (eta30 - 3*eta12)**2 + (3*eta21 - eta03)**2
phi4 = (eta30 + eta12)**2 + (eta21 + eta03)**2
phi5 = ((eta30 - 3*eta12)*(eta30 + eta12) * ((eta30 + eta12)**2 - 3*(eta21_
↪ eta03)**2) +
        (3*eta21 - eta03)*(eta21 + eta03)*(3*(eta30 + eta12)**2 - (eta21 +_
↪ eta03)**2))
phi6 = ((eta20 - eta02)*((eta30 + eta12)**2 - (eta21 + eta03)**2) +
        4*eta11*(eta30 + eta12)*(eta21 + eta03))
phi7 = ((3*eta21 - eta03)*(eta30 + eta12) * ((eta30 + eta12)**2 - 3*(eta21_
↪ eta03)**2) -
        (eta30 - 3*eta12)*(eta21 + eta03)*(3*(eta30 + eta12)**2 - (eta21 +_
↪ eta03)**2))

return dict(
    m00=m00, cx=cx, cy=cy,
    hu=[phi1, phi2, phi3, phi4, phi5, phi6, phi7]
)

```

```

[ ]: # 1) Load sequence
frames, paths = load_sequence()
print(f"Loaded {len(frames)} frames: {paths[0]} .. {paths[-1]}")

# 2) Process consecutive pairs: abs diff -> threshold -> median -> remove small_
↪ regions
results = [] # (i, mask)
for i in range(1, len(frames)):
    diff = np.abs(frames[i] - frames[i-1])

```

```

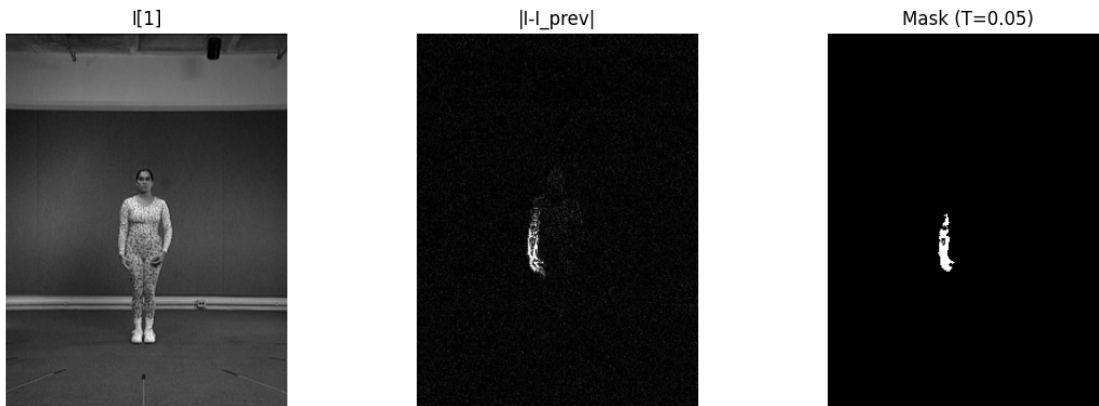
mask = (diff >= T).astype(np.uint8)
if MED_K > 1:
    mask = (median_filter2d(mask.astype(np.float32), k=MED_K) >= 0.5).
    ↪astype(np.uint8)
    mask = remove_small_regions(mask, min_size=MIN_SIZE)
    results.append((i, mask))

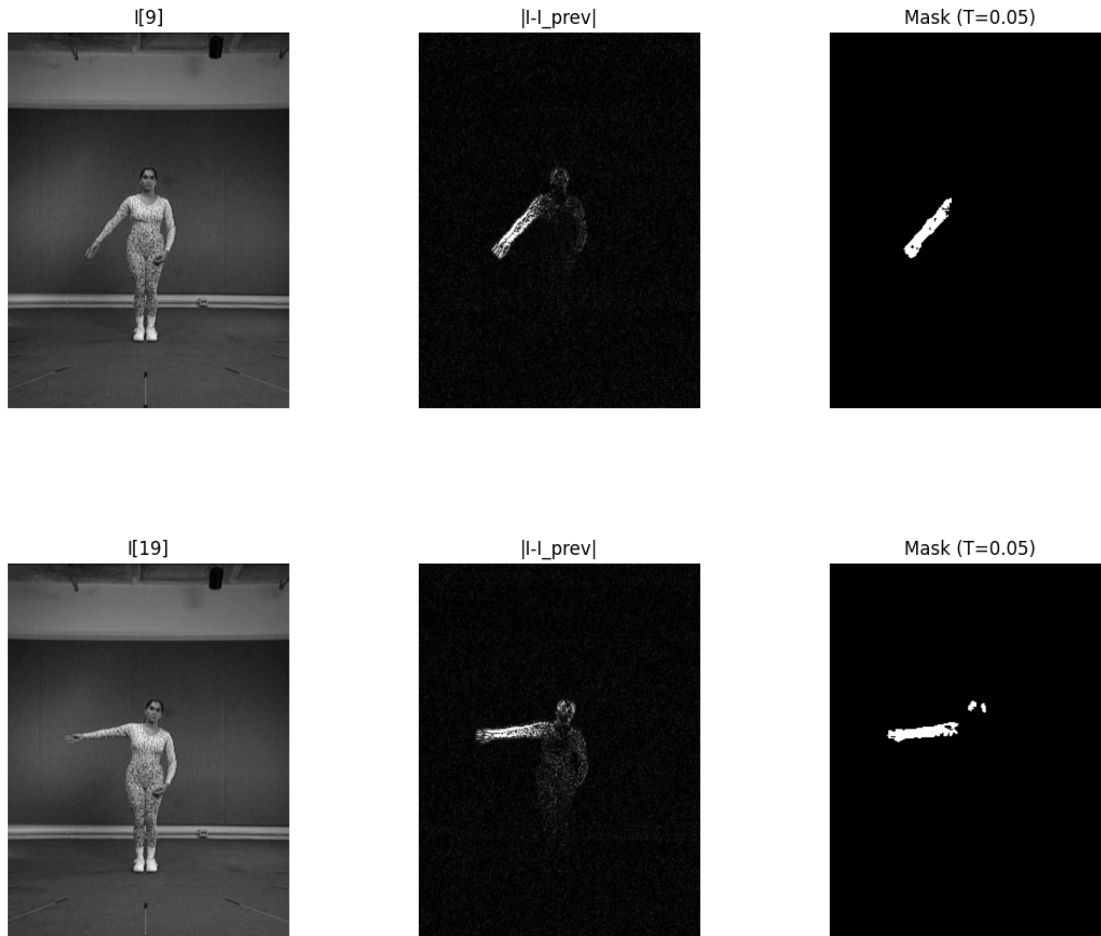
# 3) Quick visual check on a few indices
sample_is = [2, 10, 20] if len(frames) > 20 else [min(2, len(frames)-1)]
for i in sample_is:
    diff = np.abs(frames[i] - frames[i-1])
    fig, ax = plt.subplots(1, 3, figsize=(12,4))
    ax[0].imshow(frames[i-1], cmap="gray"); ax[0].set_title(f"I[{i-1}]"); ax[0].
    ↪axis("off")
    ax[1].imshow(diff, cmap="gray", vmin=0, vmax=0.2); ax[1].
    ↪set_title("|I-I_prev|"); ax[1].axis("off")
    mask = dict(results)[i]
    ax[2].imshow(mask, cmap="gray"); ax[2].set_title(f"Mask (T={T})"); ax[2].
    ↪axis("off")
    plt.tight_layout(); plt.show()

print("Done Q1.")

```

Loaded 22 frames: data/aerobic-001.bmp .. data/aerobic-022.bmp





Done Q1.

**Discussion :** Absolute frame differencing with a fixed threshold reliably highlights the moving subject, and a small median filter plus removing tiny 8-connected components cleans speckle while preserving the silhouette.

**Problem 2 :** Compute an MEI and MHI on the image sequence (using your best motion differencing approach from problem #1 above for each image pair  $i$  and  $i-1$ ), simulating the current MHI “timestamp” for each image pair using the larger of the image pair index values (i.e., use  $i$ , not  $i-1$ ). Therefore, you will have difference images from  $i=2$  to 22. The MEI/MHI duration should include all image diff results in the sequence into the final template. Use `imagesc` (Matlab) to show your results. Compute the 7 similitude moments for the final MEI and the MHI (make sure to normalize the MEI and MHI values to be between 0-1 before computing the moments using the given formula in the class notes:  $\max[0, (i-1.0)/21.0]$  for this example). [4 pts]

```
[ ]: # 1) Load frames
frames, fpaths = load_sequence()
H, W = frames[0].shape
print(f"Loaded {len(frames)} frames; computing masks & MEI/MHI...")
```

```

# 2) Build motion masks
masks = [] # list of binary masks for i=1..N-1 but we'll index by actual i (2..
            ↪N)
for i in range(1, len(frames)):
    diff = np.abs(frames[i] - frames[i-1])
    mask = (diff >= T).astype(np.uint8)
    if MED_K > 1:
        mask = (median_filter2d(mask.astype(np.float32), k=MED_K) >= 0.5).
            ↪astype(np.uint8)
        mask = remove_small_regions(mask, min_size=MIN_SIZE)
        masks.append(mask)

# 3) Compute MEI (binary OR over all masks)
MEI = np.zeros((H, W), np.uint8)
for m in masks:
    MEI |= m

# 4) Compute MHI with decay + timestamp overwrite (normalized 0..1)
# i runs 2..22 -> normalized timestamp ts = (i-1)/21 (1/21 .. 1]
MHI = np.zeros((H, W), np.float32)
for k, m in enumerate(masks, start=2):
    # decay older motion everywhere
    MHI = np.maximum(0.0, MHI - DELTA)
    # stamp current motion with normalized timestamp
    ts = max(0.0, (k - 1.0) / TAU)
    MHI[m.astype(bool)] = ts

# 5) Moments on MEI & MHI (ensure MHI in [0,1] already)
mei_mom = spatial_moments(MEI.astype(np.float32))
mhi_mom = spatial_moments(MHI)

print("\nMEI:")
print(f"  area m00 = {mei_mom['m00']:.1f}, centroid = ({mei_mom['cx']:.2f},
            ↪{mei_mom['cy']:.2f})")
print("  Hu (log10-abs):", [np.sign(h)*np.log10(abs(h)+1e-30) for h in
            ↪mei_mom['hu']])

print("\nMHI:")
print(f"  mass m00 = {mhi_mom['m00']:.1f}, centroid = ({mhi_mom['cx']:.2f},
            ↪{mhi_mom['cy']:.2f})")
print("  Hu (log10-abs):", [np.sign(h)*np.log10(abs(h)+1e-30) for h in
            ↪mhi_mom['hu']])

# 6) Quick visualization
fig, ax = plt.subplots(1, 3, figsize=(12,4))

```

```

ax[0].imshow(frames[0], cmap="gray"); ax[0].set_title("First frame"); ax[0].
    ↪axis("off")
ax[1].imshow(MEI, cmap="gray"); ax[1].set_title("MEI (binary)"); ax[1].
    ↪axis("off")
im = ax[2].imshow(MHI, cmap="gray", vmin=0, vmax=1); ax[2].set_title("MHI (0..
    ↪1)"); ax[2].axis("off")
plt.tight_layout(); plt.show()

print("Done Q2.")

```

Loaded 22 frames; computing masks & MEI/MHI...

MEI:

```

area m00 = 3284.0, centroid = (83.72, 163.20)
Hu (log10-abs): [np.float64(-0.6525978271695901),
np.float64(-2.745815789103832), np.float64(-2.0963540557258495),
np.float64(-2.76957471311622), np.float64(-5.247600866063608),
np.float64(-4.431908402158292), np.float64(5.5661535344977855)]

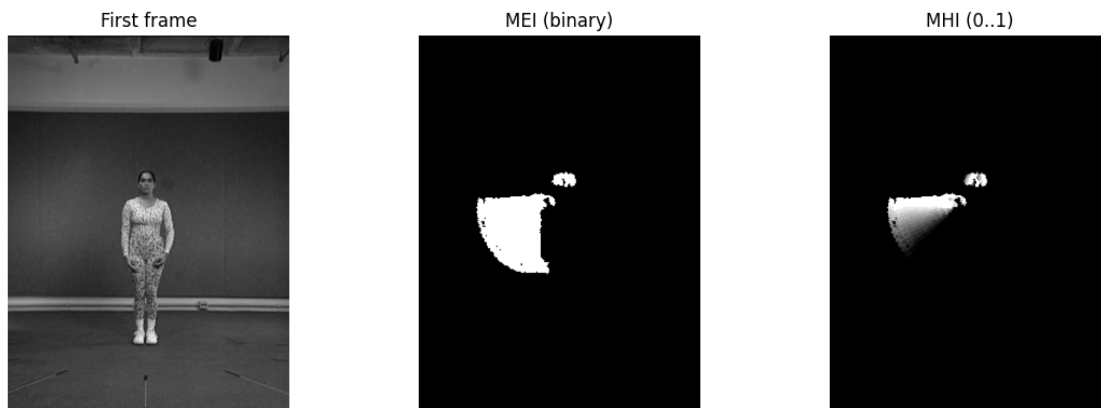
```

MHI:

```

mass m00 = 1304.1, centroid = (81.97, 145.65)
Hu (log10-abs): [np.float64(-0.31215926227854474),
np.float64(-0.83363196912822), np.float64(-1.1817623767449954),
np.float64(-1.4603369270669746), np.float64(-2.781518105438513),
np.float64(-1.879153186246054), np.float64(4.390322921860022)]

```



Done Q2.

**Discussion :** The MEI accumulates where motion occurred, while the normalized MHI (0–1) encodes when it happened, so their Hu moments emphasize overall action shape versus its recency-weighted sweep.

**Problem 3 :** Create a 101x101 image with a black (0) background and a white (255) box of size 21x21, placing the upper-left corner at pixel (row=40, col=6). Create another new box image, but

shift the box 1-pixel to the right and 1-pixel down. Compute the normal flow between the images. Use MATLAB's quiver function to draw the motion vectors on the image (call `imagesc`, then 'hold on', and lastly call `quiver`). (Make sure your gradient mask orientations/directions and the plot axes are consistent!!!) Make sure all masks are "correct" with proper scaling/normalization. Is the result what you expected? Why or why not? Comment on the flow for the 4 sides of the box and also for the 4 corners. [5 pts]

```
[25]: # 1) Build the two images (scale to [0,1] for proper normalization)
H, W = 101, 101
I1 = np.zeros((H, W), dtype=np.float32)
I2 = np.zeros((H, W), dtype=np.float32)

# Box size & positions (upper-left at (row=40, col=6), then shift +1 row, +1
↪col)
r0, c0, s = 40, 6, 21
I1[r0:r0+s, c0:c0+s] = 1.0
I2[r0+1:r0+1+s, c0+1:c0+1+s] = 1.0 # shifted down/right by 1

# 2) Spatial gradients with properly normalized Sobel masks (/8) on the 2nd
↪frame
# (slides: classic gradient masks & keep normalization factor)
# Fx = [[-1,0,1],[-2,0,2],[-1,0,1]]/8 ; Fy = [[-1,-2,-1],[0,0,0],[1,2,1]]/8
Kx = (1/8.0) * np.array([[-1, 0, 1],
                        [-2, 0, 2],
                        [-1, 0, 1]], dtype=np.float32)
Ky = (1/8.0) * np.array([[-1, -2, -1],
                        [ 0,  0,  0],
                        [ 1,  2,  1]], dtype=np.float32)

def conv2_same(img, k):
    r = k.shape[0]//2
    P = np.pad(img, r, mode="reflect")
    out = np.empty_like(img, dtype=np.float32)
    for y in range(img.shape[0]):
        for x in range(img.shape[1]):
            out[y, x] = np.sum(P[y:y+2*r+1, x:x+2*r+1] * k, dtype=np.float32)
    return out

Ix = conv2_same(I2, Kx)
Iy = conv2_same(I2, Ky)

# 3) Temporal derivative (It = I2 - I1), dt=1 frame
It = I2 - I1

# 4) Normal flow: solve Ix*u + Iy*v + It = 0 for the component along I
# v_n = -(It / (Ix^2 + Iy^2)) * [Ix, Iy]
eps = 1e-12
```



```

den = Ix*Ix + Iy*Iy + eps
U   = -It * Ix / den
V   = -It * Iy / den

# Keep vectors only where spatial gradient is significant (on edges)
mag = np.sqrt(Ix*Ix + Iy*Iy)
thr = 0.2 * mag.max() # show clear edge-only vectors
mask = mag >= thr
U[~mask] = 0.0
V[~mask] = 0.0

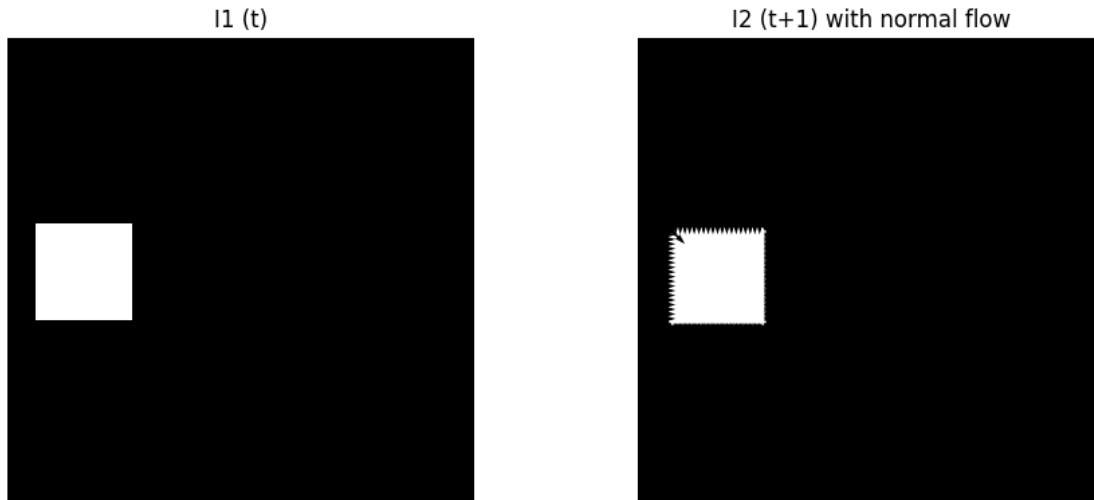
# 5) Plot: image then quiver (MATLAB: imagesc; hold on; quiver)
Y, X = np.mgrid[0:H, 0:W]
fig, ax = plt.subplots(1, 2, figsize=(10, 4))

ax[0].imshow(I1, cmap="gray", vmin=0, vmax=1, origin="upper")
ax[0].set_title("I1 (t)")
ax[0].axis("off")

ax[1].imshow(I2, cmap="gray", vmin=0, vmax=1, origin="upper")
ax[1].quiver(X[mask], Y[mask], U[mask], V[mask],
             angles="xy", scale_units="xy", scale=1)
ax[1].set_title("I2 (t+1) with normal flow")
ax[1].invert_yaxis() # ensure row+ points downward (image coords)
ax[1].set_aspect("equal")
ax[1].set_xlim(0, W-1); ax[1].set_ylim(H-1, 0)
ax[1].axis("off")

plt.tight_layout()
plt.show()

```



**Discussion :** Yes. Normal flow only picks up motion straight out from edges, so arrows go sideways on vertical edges and up/down on horizontal edges; corners look messy/weak, and inside the box there are no arrows because there's no texture to measure.